

Distributed Canary Deployment – Design Plan

Selin Demirtürk, Roberto Villafuerte, Jan Thurner

November 14, 2025

Contents

| | | |
|----------|--|----------|
| 1 | Introduction | 1 |
| 2 | Topics and Techniques | 2 |
| 2.1 | Global Routing State | 2 |
| 2.2 | Consistency and Consensus (2PC) | 2 |
| 2.3 | Operational Parameters (SLOs, Timeouts, Ports) | 3 |
| 2.4 | Failure Modes & Recovery | 3 |
| 2.4.1 | Coordinator crash during PREPARE | 3 |
| 2.4.2 | Participant crash before vote | 3 |
| 2.4.3 | Participant crash after PREPARED | 3 |
| 2.4.4 | Network partition | 3 |
| 2.5 | Persistence and Recovery | 3 |
| 2.6 | Scalability Considerations | 3 |
| 3 | Nodes and their Roles | 4 |
| 3.1 | Common | 4 |
| 3.2 | Coordinator | 4 |
| 3.3 | Participants | 4 |
| 4 | Messaging | 5 |
| 4.1 | Control Plane (peer ↔ peer) | 5 |
| 4.2 | Data Plane (client → server) | 5 |
| 4.3 | Message Summary | 5 |

1 Introduction

We design a distributed Canary Model Rollout Orchestrator for ML services at global scale. Production currently serves model v1. A new model v2 is introduced gradually (e.g., 5% → 10% → 50% → 100%) only if health is good. A tiny global routing state (`{model_id -> weight, version, status}`) is agreed atomically across all serving nodes so that every region flips to the same configuration at the same logical moment. If latency/error SLOs degrade, the system aborts the change or rolls back to the last good version. The full-scale vision can be used for multiple regions with hundreds of nodes per region, each region makes local routing decision using the same protocol.

Top-level Architecture




Figure 1: Three-node prototype (scales to N). Each node talks to the other two.

2 Topics and Techniques

We implement a distributed *canary rollout orchestrator* where all serving nodes agree on a tiny, versioned `RoutingState` ($\{\text{model} \rightarrow \text{weight}, \text{version}, \text{status}\}$). A change (e.g., 5% \rightarrow 10%) is published via **Two-Phase Commit (2PC)**: during PREPARE each node first *stages* the canary artifact (verify checksum, warm-load the model), runs quick *smoke tests* and local health gates, and persists PREPARED. If every node votes COMMIT, the coordinator broadcasts GLOBAL_COMMIT; each node atomically flips its in-memory router to the new weights, starts serving per those weights, and persists COMMITTED. Otherwise, GLOBAL_ABORT keeps the previous version. Messages are idempotent (carry `txid`, `version`); state transitions are durably logged and recovered by replay/peer queries, preventing mixed views after crashes or partitions.

2.1 Global Routing State

We maintain a single source of truth replicated on all nodes:

Listing 1: `RoutingState` (persisted and agreed cluster-wide).

```
{
  "version": 18,                                // monotonically increasing
  "stable_model_id": "v1",
  "canary_model_id": "v2",
  "weights": { "v1": 0.95, "v2": 0.05 }, // sum = 1.0
  "status": "PREPARED|COMMITTED|ABORTED",
  "txid": "promote-2025-11-12T10:22:31Z"
}
```

2.2 Consistency and Consensus (2PC)

We use **Two-Phase Commit** to make state changes atomic:

1. **Phase 1: PREPARE.** Coordinator proposes next version (e.g., 5%→10%). Each participant: warm-loads model, checks health gates, persists PREPARED <version>, then votes COMMIT or ABORT.
2. **Phase 2: DECISION.** If all votes are COMMIT, coordinator broadcasts COMMIT <version>; else broadcasts ABORT <version>. Participants atomically apply the decision and persist it.

Stretch goal: Replace 2PC with Raft over a change log to allow majority progress under certain failures.

2.3 Operational Parameters (SLOs, Timeouts, Ports)

- **Health gates (window = min{60 s, 2000 reqs}):** $p95_{\text{lat}} \leq 1.2 \times p95_{\text{baseline}}$, error_rate $\leq 0.5\%$.
- **2PC timeouts:** PREPARE reply timeout = 2 s; decision timeout = 1 s; backoff jitter 100–300 ms; max retries = 3.
- **Ports:** Node A: 50051; Node B: 50052; Node C: 50053.
- **Snapshots:** every 5 committed versions; retain last 50 log entries.

2.4 Failure Modes & Recovery

2.4.1 Coordinator crash during PREPARE

Participants time out, elect a temporary resolver (highest node id), query peers for decision: commit if any peer has COMMITTED <v>, else ABORT <v>. Guarantees no mixed views.

2.4.2 Participant crash before vote

Coordinator times out and decides ABORT. Crashed node replays log on restart and serves last COMMITTED version.

2.4.3 Participant crash after PREPARED

Coordinator may still ABORT; on restart, node finds PREPARED <v> without decision, queries peers, applies final decision.

2.4.4 Network partition

2PC requires unanimous votes; promotion pauses or aborts under partition. Already committed versions remain; nodes catch up via heartbeat upon heal.

2.5 Persistence and Recovery

Each node keeps an append-only JSON log of state transitions and periodically snapshots the last COMMITTED state. On restart, a node replays its log to reconstruct state. If it finds a highest PREPARED version without decision, it queries peers: commit if any peer has COMMITTED, otherwise abort. Heartbeats exchange {node_id, last_committed_version, digest} for catch-up.

2.6 Scalability Considerations

The control plane manipulates a tiny object; its cost grows slowly with node count. The data plane scales horizontally by adding serving replicas. Global deployment can be organized per-region with the same protocol, then orchestrated region-by-region for staged rollouts.

3 Nodes and their Roles

3.1 Common

- **Serving:** /predict routes requests by local weights of the latest COMMITTED version (optional sticky sessions).
- **State Manager:** Holds in-memory RoutingState; persists PREPARED/COMMITTED/ABORTED; snapshots; replay on startup.
- **Health Watcher:** Maintains rolling metrics; exposes /health/snapshot.
- **Peer Comms:** Sends/receives heartbeats; requests missing decisions to catch up.

3.2 Coordinator

- Proposes next version (e.g., 5%→10%) and runs 2PC.
- Applies timeouts; if a participant is unresponsive, decides ABORT (safety first).
- Triggers rollback (proposes prior stable version) if SLO breach is detected.

3.3 Participants

- On PREPARE: warm-load, check gates, persist PREPARED, vote.
- On decision: atomically swap weights; persist COMMITTED/ABORTED.

4 Messaging

4.1 Control Plane (peer ↔ peer)

Listing 2: Core control messages.

```

message RoutingState {
    int64 version;
    string stable_model_id;
    string canary_model_id;
    // weights serialized as "model_id:share" pairs
    repeated string weights_k;    // e.g., "v1"
    repeated double weights_v;   // e.g., 0.95
    string status;               // PREPARED|COMMITTED|ABORTED
    string txid;
}

message PrepareReq { int64 version; RoutingState state; string txid; }
message PrepareResp { string txid; enum Vote { COMMIT=0; ABORT=1; } vote;
                      string reason; }

message Decision { int64 version; string txid; enum Kind { COMMIT=0; ABORT=1; }
                     kind; }

message Ack { string txid; bool ok; }

message HealthSnap { string node_id; double p95_ms; double err_rate; int64 n;
                      string window_id; }

message Heartbeat { string node_id; int64 last_committed_version; bytes digest; }

```

4.2 Data Plane (client → server)

Listing 3: HTTP endpoints (example).

```

POST /predict           {input...}      ->
routed to v1 or v2 by local weights (current COMMITTED version)

GET /routing/state     -> {version, weights, model_ids, status}

GET /health/snapshot   -> {p95_ms, err_rate, n, window_id}

```

4.3 Message Summary

| Name | Direction | Purpose / Key Fields |
|-------------|----------------|--|
| PrepareReq | Coord. → Part. | Propose new version; includes RoutingState, version, txid. |
| PrepareResp | Part. → Coord. | Vote COMMIT/ABORT; reason. |
| Decision | Coord. → Part. | Final decision COMMIT/ABORT for version. |
| Ack | Part. → Coord. | Acknowledge decision applied. |
| HealthSnap | Peer ↔ Peer | p95, error rate, window id, sample size n. |
| Heartbeat | Peer ↔ Peer | last_committed_version, digest. |




Figure 2: Canary rollout communication: coordinator issues PREPARE; participants stage the canary, run health gates, persist PREPARED, and vote. On unanimous commit the coordinator sends GLOBAL_COMMIT (nodes persist COMMITTED and atomically flip routing); otherwise GLOBAL_ABORT. All decisions are acknowledged and durably recorded in the local RoutingState store.